

# Lab 4: TCP Congestion Control

50.012 Networks

Hand-out: Nov 6

Hand-in: Nov 18 (Monday 11:59pm)

## 1 Objectives

- Experiment with TCP congestion control and study the effect of router buffer sizes.
- Goal: 1) Understand TCP's congestion control mechanism. 2) Get familiar with network emulation. 3) Understand why larger buffers can be disadvantageous.
- Part of this exercise is based on the Stanford CS244 class lab1.

## 2 Setup

### 2.1 Lab4 setup

- Download the lab4.zip from eDimension and unpack to some local folder, e.g. ~/lab4/.
- Change into the lab4 folder, and open up another terminal tab (e.g., with CTRL+SHIFT+T).

### 2.2 mininet

- This exercise assumes that you have a running mininet installation on Ubuntu 14.04 (please refer to lab 4 preparation guide in eDimension for that).

#### 1. What is Mininet?

- Mininet is a network emulator (mostly written in Python), that allows you to define network topologies and to connect emulated hosts.
- The emulated hosts run on your OS, and can execute your client/server applications easily.
- You can also easily change link parameters such a packet loss rate, delay, and bandwidth.
- Mininet is mainly intended to experiment with Software Defined Networks and Open-Flow, but is also convenient for this lab session.
- More details at <http://mininet.org>, in particular <http://mininet.org/walkthrough/>
- If you like videos: some basic things of this sheet are also demo'ed in <https://www.youtube.com/watch?v=jmlgXaocwiE>

## 2. Installation and first test

- Start up mininet:

```
sudo ./run.sh
```

- Get familiar with mininet commands using mininet's help system:

```
mininet> help [topic]
```

- Use the following mininet command to spawn two terminals, belonging to two nodes in the emulated network.

```
mininet> xterm h1 h2
```


- Using these terminals, you can run applications or perform ping operations on the emulated hosts
  - To close running commands like `ping -c 100`, try CTRL-C
- Find out the IP addresses of h1 and h2
- Note that `mininet` allows to run a command from its command prompt as if the command is run from any virtual host. For example, if we want to ping h2 from h1 we can run:

```
mininet> h1 ping h2
```

- To leave/close Mininet, try CTRL-D on the main command line


## 3 Queue Length and TCP

In this exercise we will study the dynamics of TCP in home networks. Take a look at the figure below which shows a "typical" home network with a "Home Router" connected to an end host. The "Home Router" is connected via Cable or DSL to a Headend router at the Internet access provider's office. We are going to study what happens when we download data from a remote server to the End Host in this home network.



Problem overview

In a real network it's hard to measure the TCP congestion window  $cwnd$  (because it's private to the Server) and the buffer occupancy (because it's private to the router). Luckily, we have mininet to help allow us to get these values easily.



### 3.1 Simple TCP connection (wget)

- Start up mininet

```
sudo ./run.sh
```

- In the running mininet emulation, test how long it takes to download a webpage on h1 from h2.
- wget's visualization may be a bit confusing. You can find the total time by subtracting start time from end time, or by looking at the time at end of last part line.

```
mininet> h2 wget h1
```

- We discussed the TCP congestion control, what is your expectation on how the cwnd changes during this HTTP session?

### 3.2 With parallel load (iperf)

- To see how the dynamics of a long flow differs from a short flow (which never leaves slow-start), we are going to repeat Part 2 for a “streaming video flow”. Instead of actually watching videos on your machine, we are going to set up a long-lived high speed TCP connection instead, to emulate a long-lived video flow.
- You can generate long flows using the `iperf` command. `iperf` is a tool for performing network throughput measurements. It can test either TCP or UDP throughput. To perform an `iperf` test the user must establish both a server (to discard traffic) and a client (to generate traffic). We have wrapped `iperf` in a script that you can run as follows:

```
mininet> h1 ./iperf.sh
```

- You can see the throughput of TCP flow from h1 to h2 by running:

```
mininet> h2 tail -f ./iperf-recv.txt
```

- You can quit viewing throughput by pressing CTRL-C.
- Think about the congestion window of the TCP stream again. How will it look for a long-established connection?
- You can also use ping to observe RTT while iperf is running

```
mininet> h1 ping -c 100 h2
```

#### The Impact on the Short Flow

- To see how our long-lived iperf flow affects our web page download, download the webpage again - while iperf is running. Observe how long it takes.

```
mininet> h2 wget h1
```

- Why does the web page take so much longer to download?

### 3.3 Measuring the real cwnd and buffer occupancy values.

We provided a script that lets you measure cwnd and buffer occupancy values (in terms of values for cwnd and buffer occupancy). We are going to re-run a couple of the experiments and plot the real values.

#### 1. Restart Mininet

- Stop and restart Mininet and the monitor script, then re-run the above experiment as follows.

```
mininet> exit  
sudo ./run.sh
```

#### 2. Monitor TCP CWND and Buffer Occupancy in Mininet

- In your other terminal tab, type the following giving a name for your experiment.

```
./monitor.sh <EXP_NAME>
```

- In your first tab with the running mininet, start iperf again

```
mininet> h1 ./iperf.sh
```

- (wait for 70 seconds ...)

```
mininet> h2 wget h1
```

- Wait for the wget to complete, then stop the python monitor script followed by the instructions on the screen. The cwnd values are saved in <EXP\_NAME>\_tcpprobe.txt and the buffer occupancy in <EXP\_NAME>\_sw0-qlen.txt.

#### 3. Plot CWND and Queue Occupancy

- Plot the TCP cwnd and queue occupancy from the output file

```
./plot_figures.sh <EXP_NAME>
```

- Adjust command line parameters to generate the figure you want.

- The script will also host a webserver on the machine and you can use the url the script provided to access to your figures.
- Note: the figures will sometimes interpolate where it is not appropriate. If you see a slow decrease of cwnd, then that probably is an artefact of such interpolation. Ignore those segments of the figures.
- If you are unable to see the cwnd, ensure you run wget after you started the monitor.sh script.

### 3.4 Smaller Buffer: from 100 packets to 20 packets

- Stop any running Mininet. This time we will make the buffers 20 packets long instead using:

```
sudo ./run-minq.sh
```

- Repeat the earlier simple tests

```
mininet> h2 wget h1
mininet> h1 ping -c 10 h2
```

- Did the results change?
- Close mininet, restart the experiment and start monitoring again

```
sudo ./run-minq.sh
sudo ./monitor.sh <EXP_NAME>
```

- Re-run the experiment:

```
mininet> h1 ./iperf.sh
mininet> h1 ping -c 30 h2
mininet> h2 wget h1
```

- What do you think the cwnd and queue occupancy will be like in this case?
- Plot the figure for cwnd and queue occupancy, this time using:

```
./plot_figures_minq.sh <EXP_NAME>
```

- Then again, use the url to see your figures. Why does reducing the queue size reduce the download time for wget?

## 4 What to Hand in

### 4.1 eDimension submission:

- Please provide a writeup (PDF format with your name) that includes the following information:
  - What is the normal time required to download the webpage on h1 from h2?
  - What was your initial expectation for the congestion window size over time?

- After starting iperf on h1, did you observe something interesting in the ping RTT?
  - After starting iperf on h1, why does the web page take so much longer to download?
  - Please provide the figures for the first experiment (with qlen 100)
    - \* Please comment on what you can see in the figures
  - Please provide the figures for the second experiment (with qlen 20)
    - \* Please comment on what you can see in the figures, and what is different (and why)
- You could work in a pair to conduct the experiments, but please summarize your observations and submit the writeup individually.